

# CSC 665: Artificial Intelligence

## Games: Adversarial Search

Instructor: Pooyan Fazli  
San Francisco State University

# How to Succeed in This Course

---

- Read all the slides
- Read all the required reading material
- Read the “Review Materials” on iLearn
- Concepts in this course take some time to sink in: be careful not to fall behind.  
Cramming will NOT do it
- Be active in lectures and on iLearn.
- Study in groups
- Ask us if you have the slightest doubt

# Course Information: Feedback

- Please give feedback (positive or negative) as often as and as early as you can.

## CSC 665: Anonymous Feedback

Name (Optional)

Email Address (Optional)

Any Feedback on CSC 665?

Submit

*Never submit passwords through Google Forms.*

<https://forms.gle/jeuK3BNmGW7vRaGW7>

# Game Playing State-of-the-Art

## ■ Checkers:

- 1950: First computer player.
- 1959: Samuel's self-taught program.
- 1994: First computer world champion: Chinook ended 40-year-reign of human champion Marion Tinsley.
- 2007: Checkers solved! Endgame database of 39 trillion states

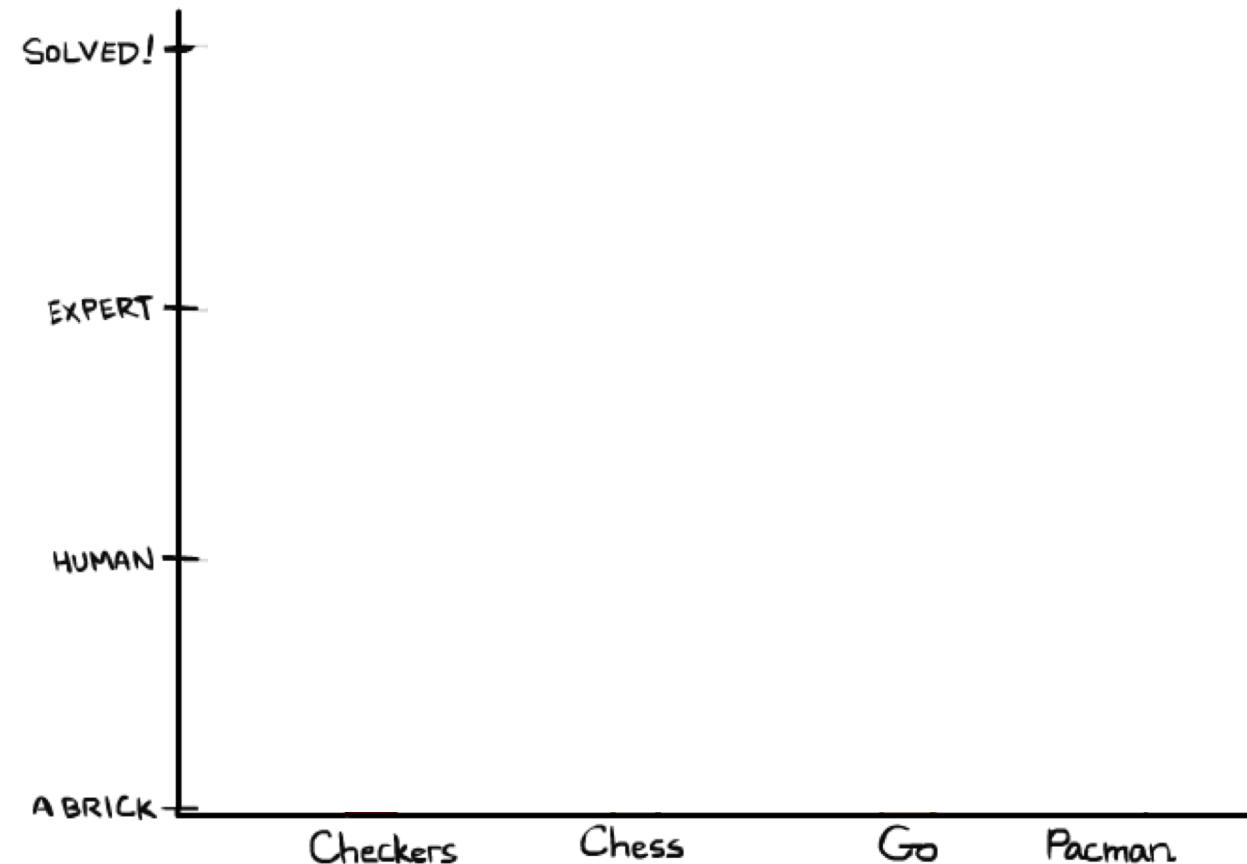
## ■ Chess:

- 1945-1960: Zuse, Wiener, Shannon, Turing, Newell & Simon, McCarthy.
- 1960s onward: gradual improvement
- 1997: special-purpose chess machine Deep Blue defeats human champion Gary Kasparov in a six-game match. Deep Blue examined 200M positions per second and extended some lines of search up to 40 ply. Current programs running on a PC rate > 3200 (vs 2870 for Magnus Carlsen).

## ■ Go:

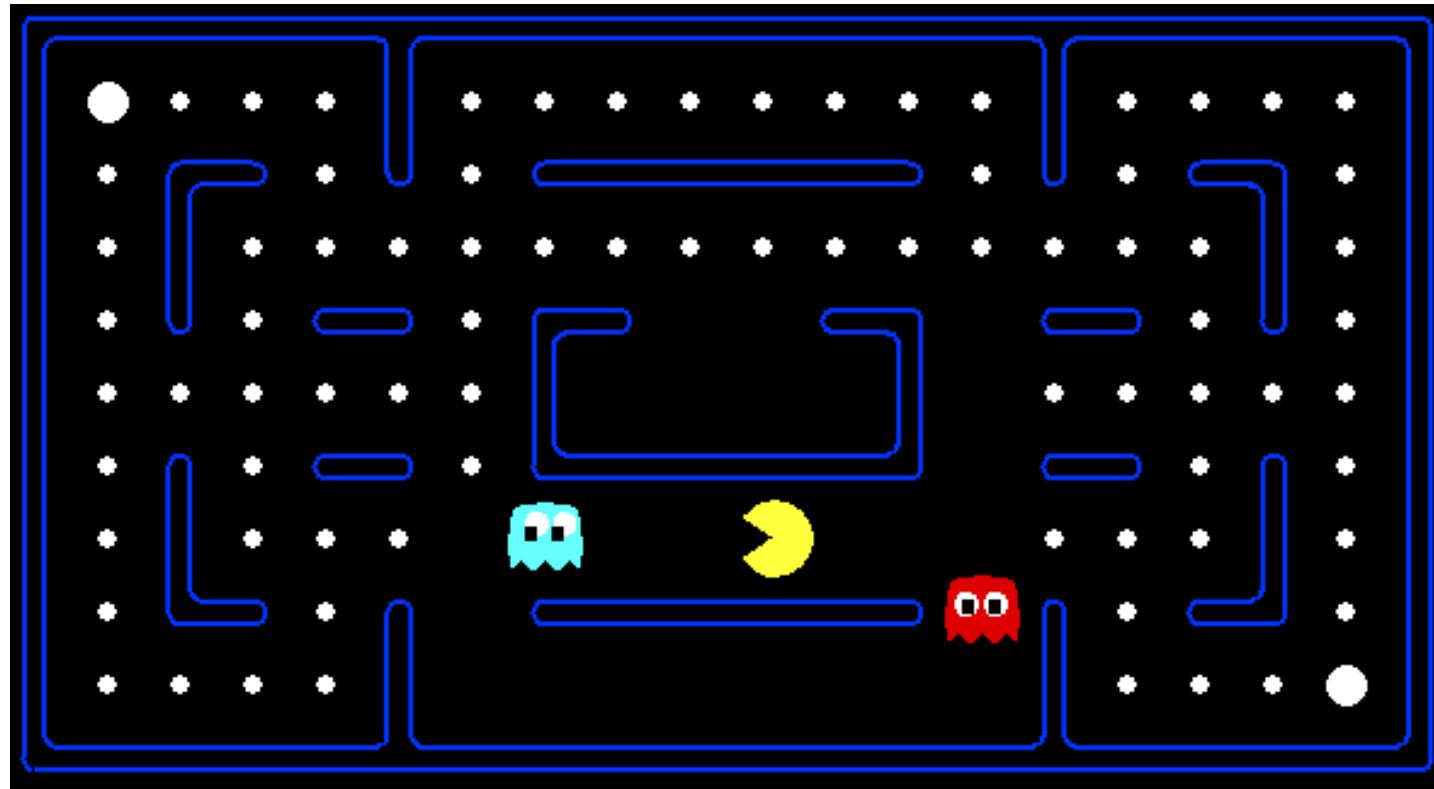
- 1968: Zobrist's program plays legal Go, barely ( $b>200!$ )
- 2005-2014: Monte Carlo tree search enables rapid advances: current programs beat strong amateurs, and professionals
- 2016: Google AI Beat World Go Champion

## ■ Pacman



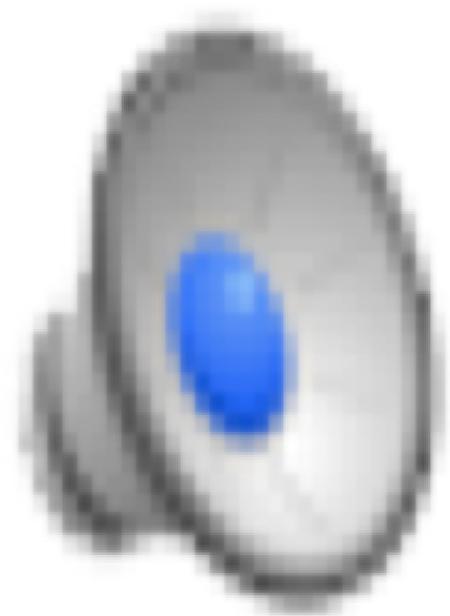
# Behavior from Computation

---



# Video of Demo Mystery Pacman

---



# Adversarial Games

---

# Types of Games

---

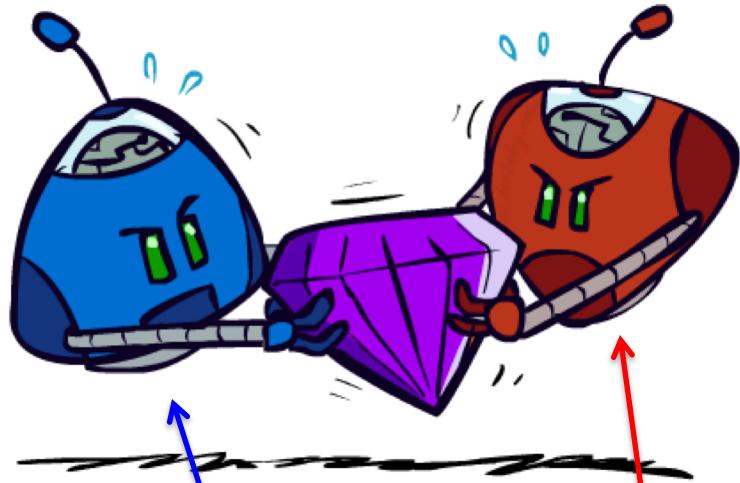
- Many different kinds of games!
- Axes:
  - Deterministic or stochastic?
  - Perfect information (fully observable)?
  - One, two, or more players?
  - Turn-taking or simultaneous?
  - Zero sum?
- Want algorithms for calculating a **strategy (policy)** which recommends a move from each state

# Deterministic Games

---

- Many possible formalizations, one is:
  - States:  $S$  (start at  $s_0$ )
  - Players:  $P=\{1\dots N\}$  (usually take turns)
  - Actions:  $A$  (may depend on player / state)
  - Transition Function:  $S \times A \rightarrow S'$
  - Terminal Test:  $S \rightarrow \{t, f\}$
  - Terminal Utilities:  $S \times P \rightarrow R$
- Solution for a player is a **policy**:  $S \rightarrow A$

# Zero-Sum Games



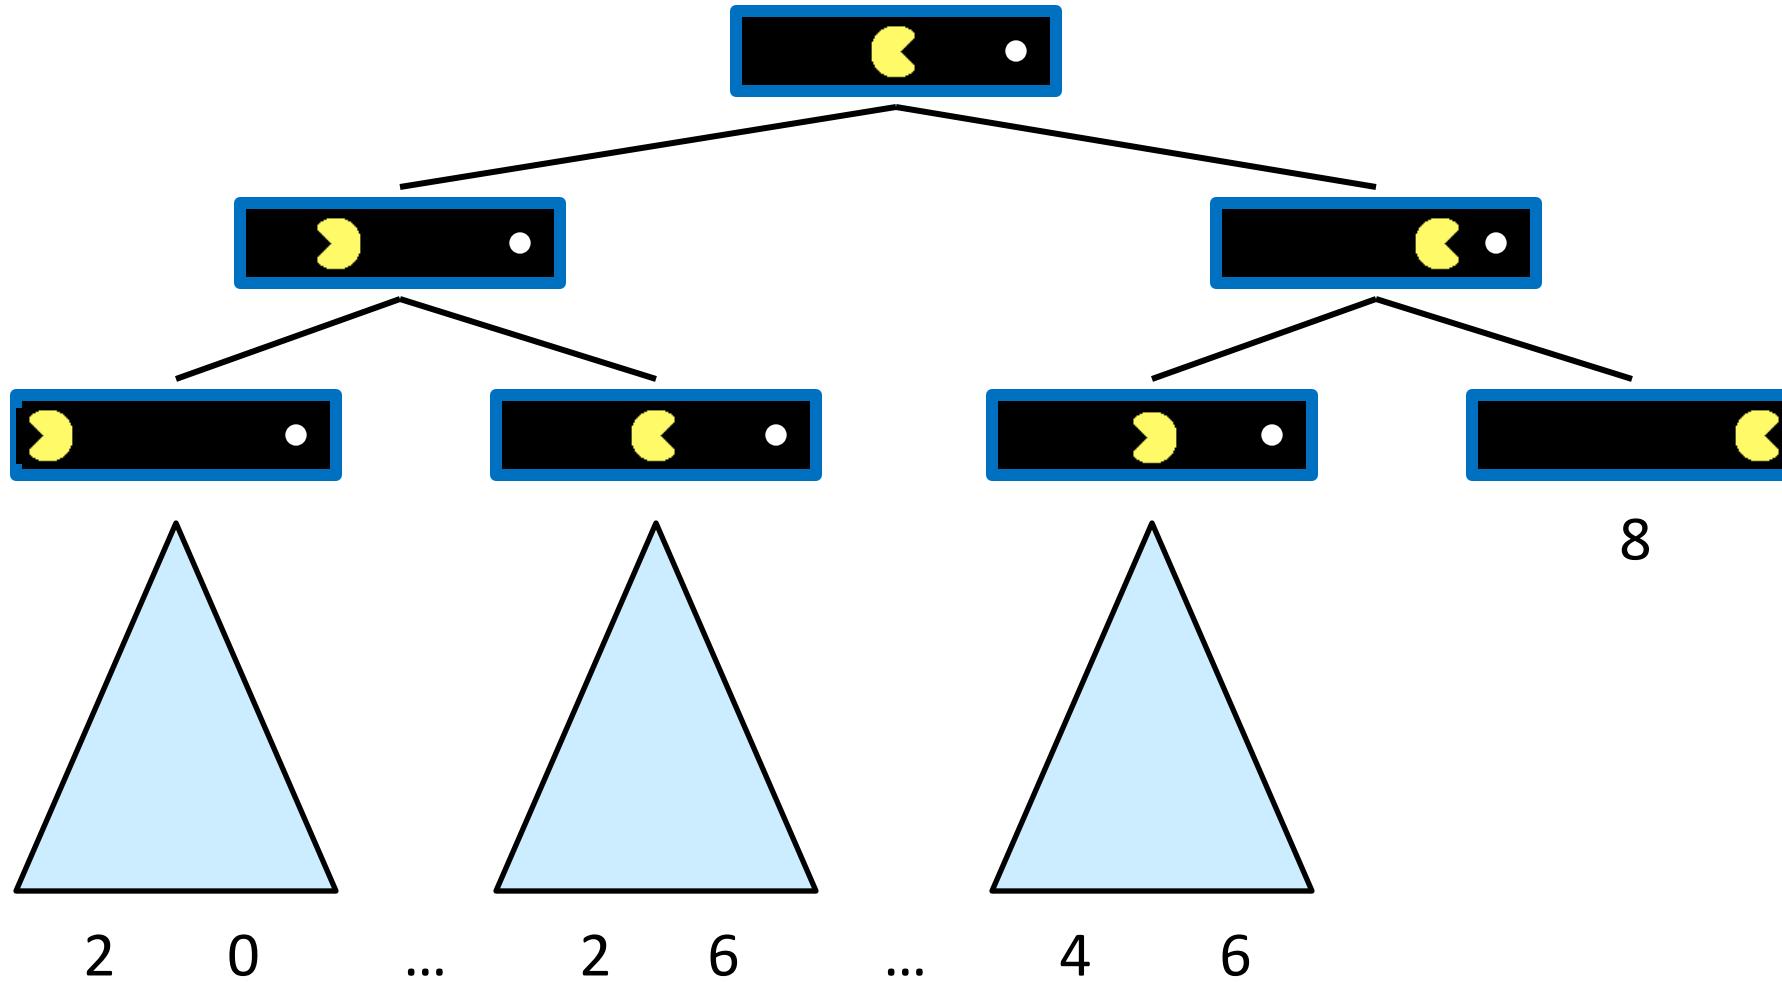
- Zero-Sum Games
  - Agents have **opposite** utilities
  - Pure competition:
    - One **maximizes**, the other **minimizes**

- General Games
  - Agents have **independent** utilities
  - Cooperation, indifference, competition and more are all possible

# Adversarial Search

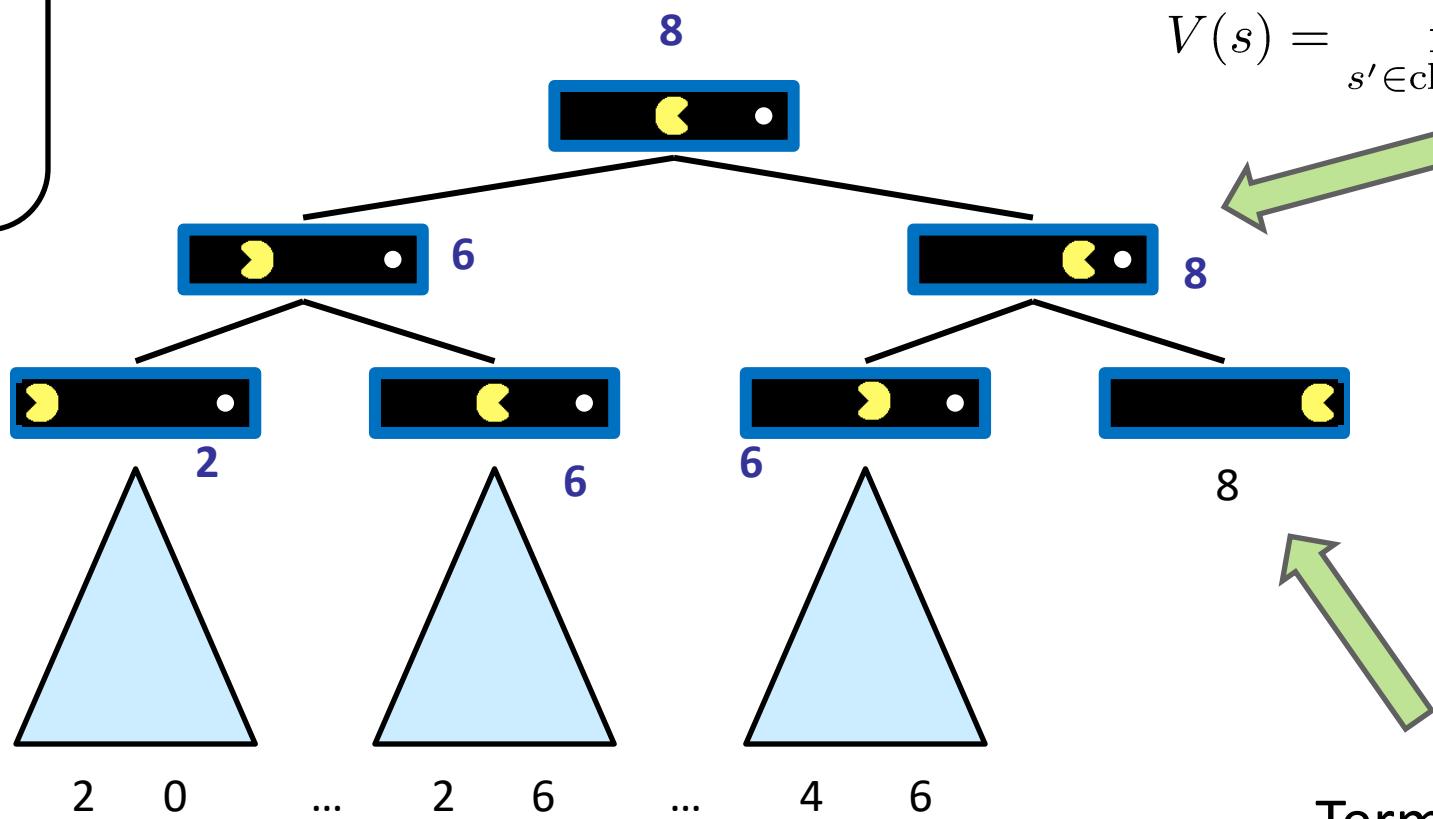
---

# Single-Agent Trees



# Value of a State

Value of a state:  
The best achievable  
outcome (utility)  
from that state



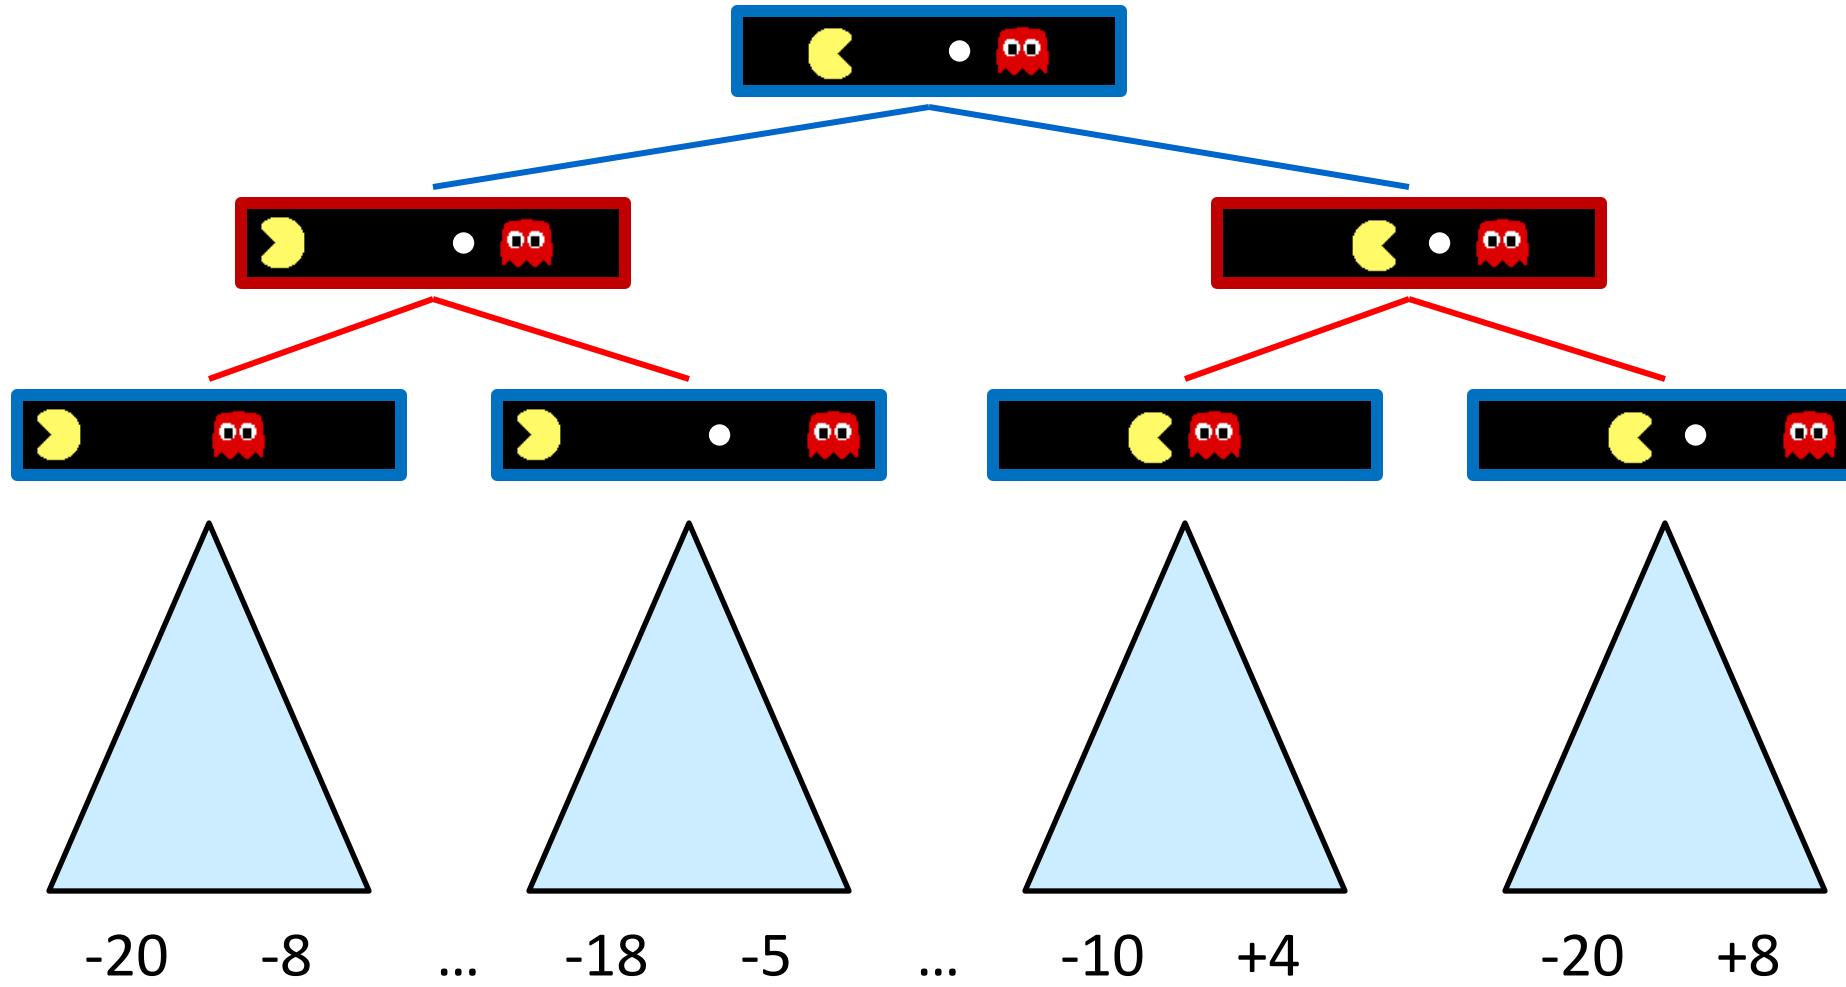
Non-Terminal States:

$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

Terminal States:

$$V(s) = \text{known}$$

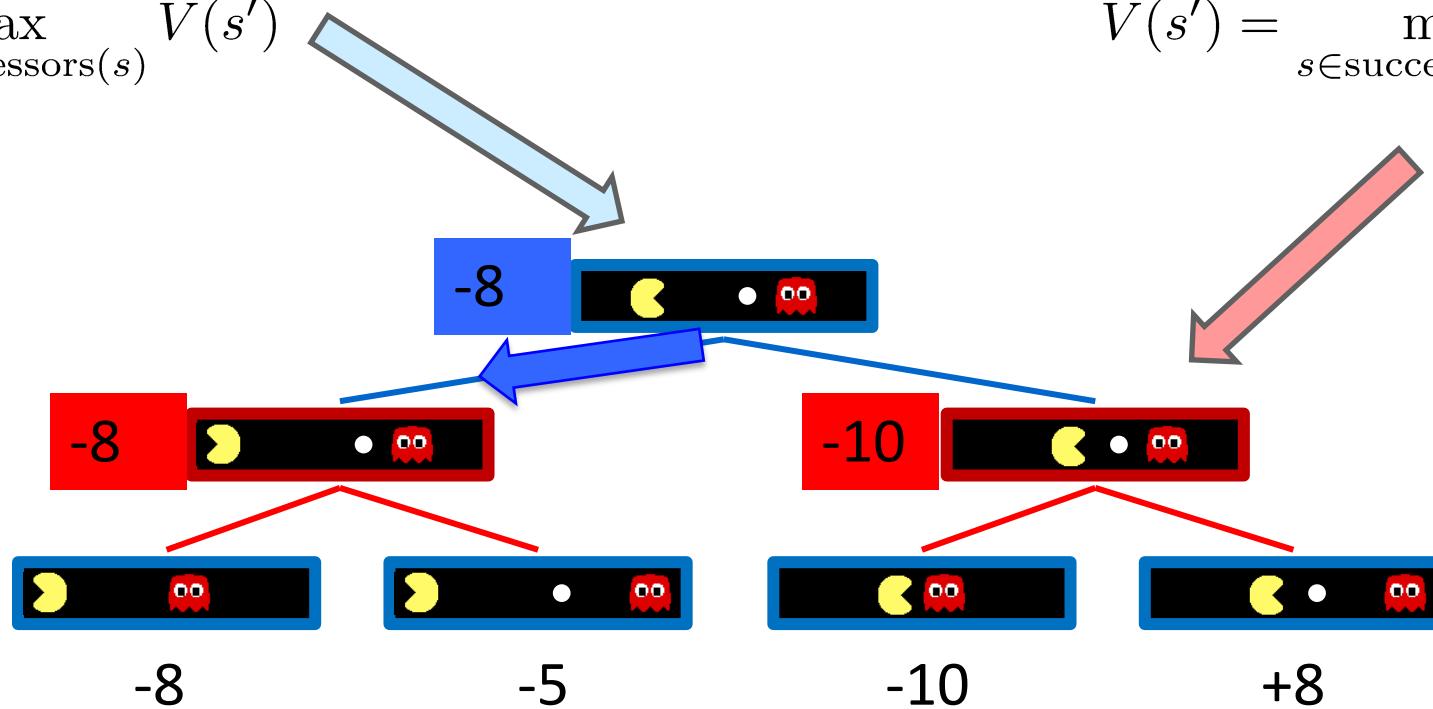
# Adversarial Game Trees



# Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$



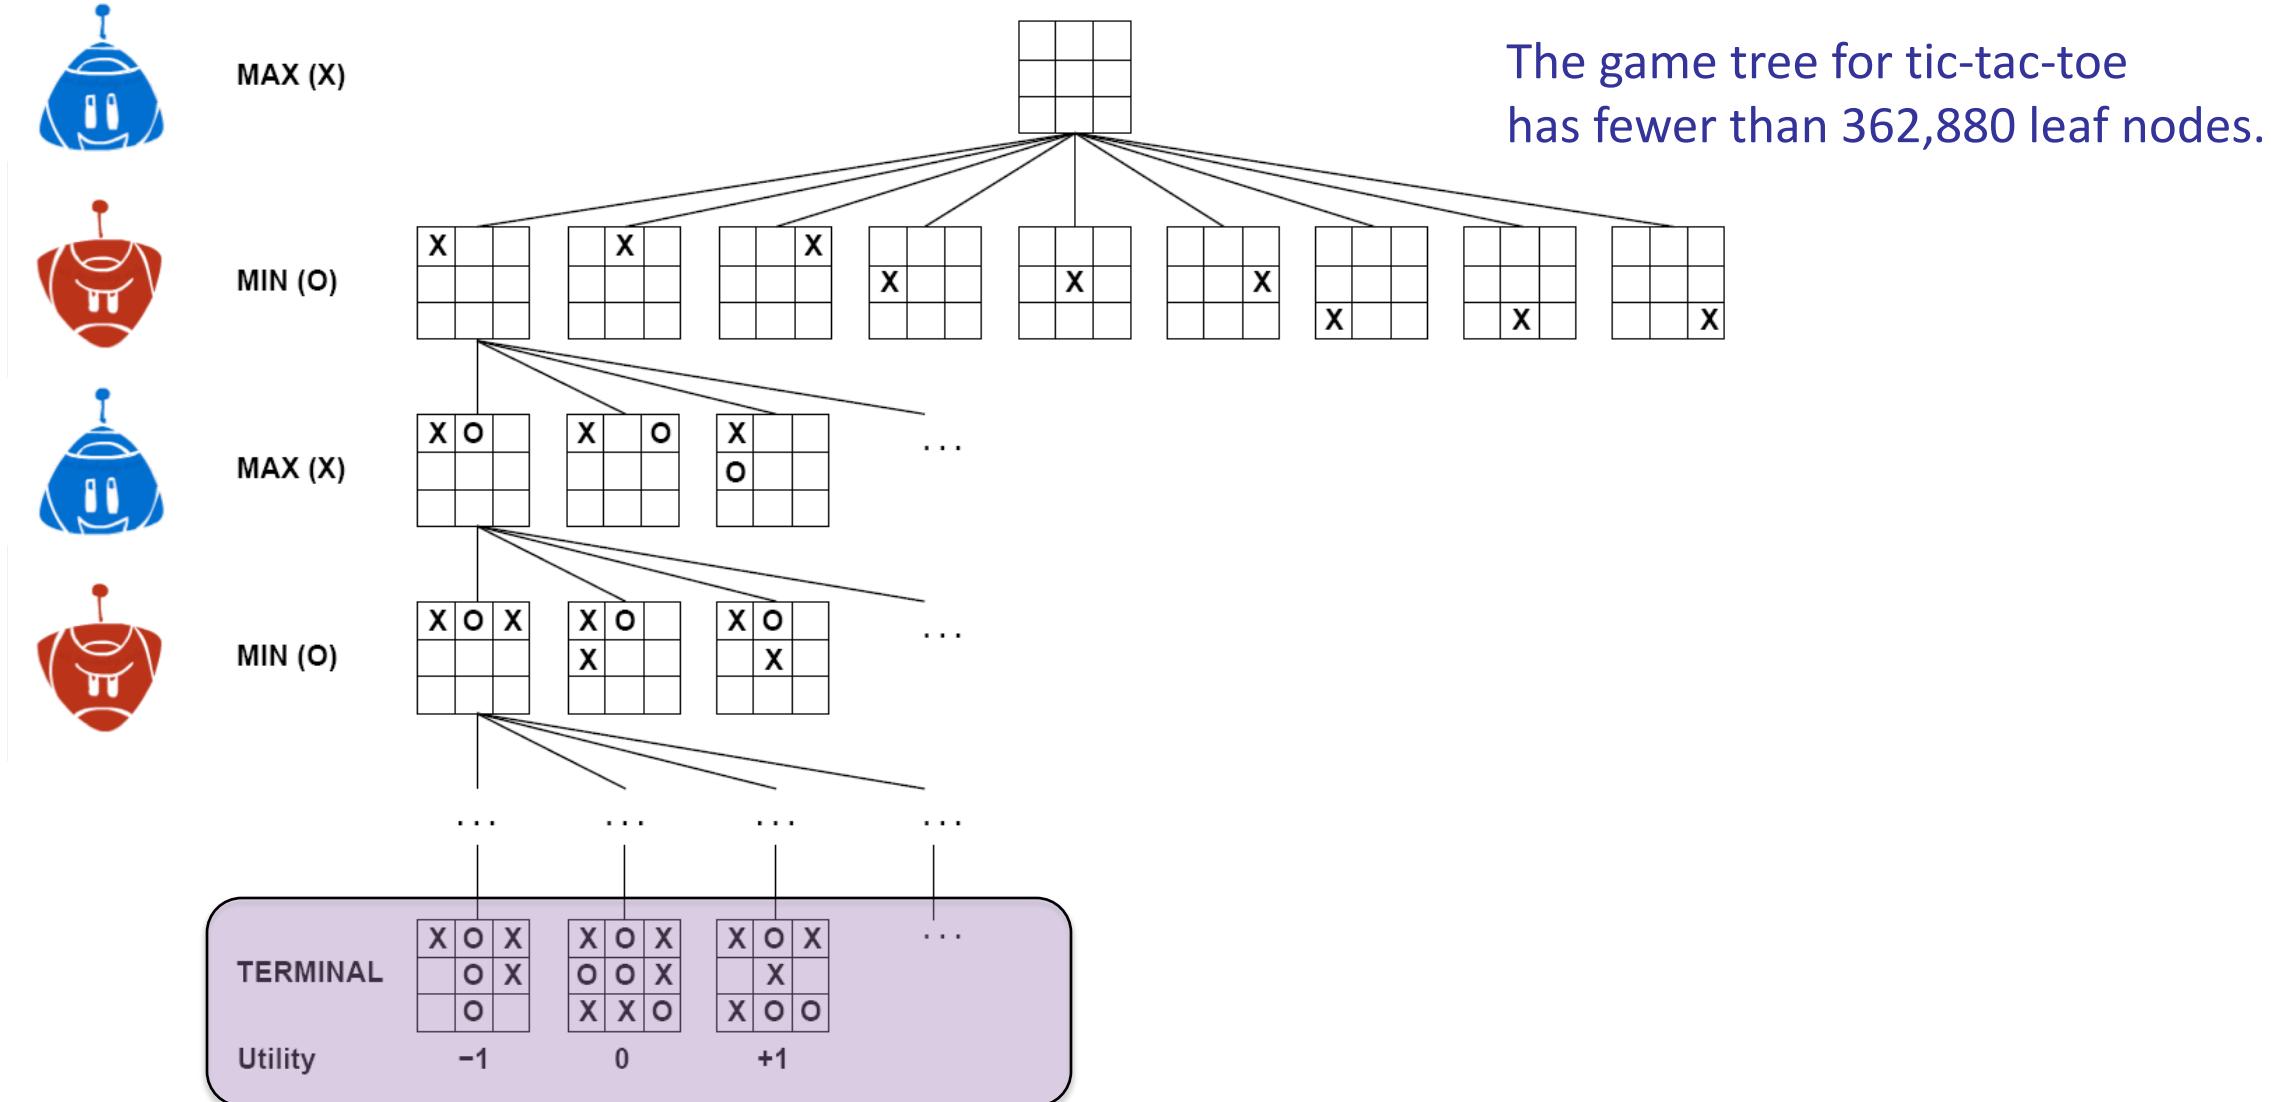
States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Terminal States:

$$V(s) = \text{known}$$

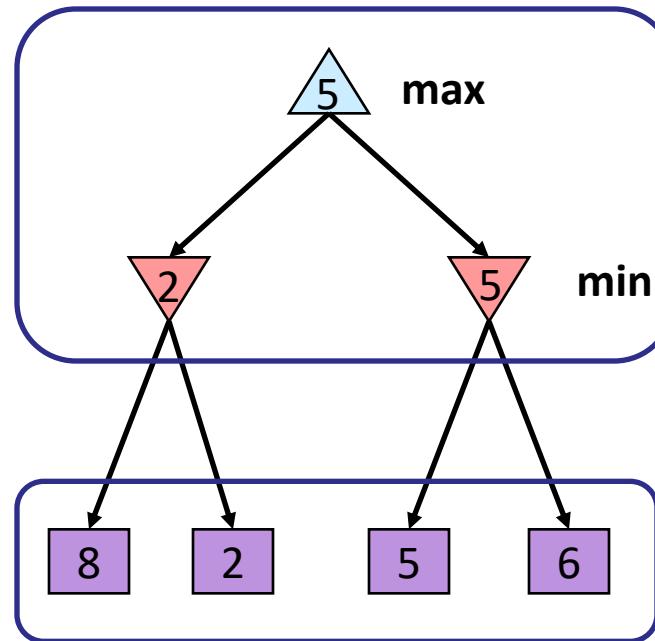
# Tic-Tac-Toe Game Tree



# Adversarial Search (Minimax)

- Deterministic, zero-sum games:
  - Tic-tac-toe, chess, checkers
  - One player maximizes result
  - The other minimizes result
- Minimax search:
  - A state-space search tree
  - Players alternate turns
  - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary

Minimax values:  
computed recursively



Terminal values:  
part of the game

# Minimax Implementation

```
def max-value(state):  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```



```
def min-value(state):  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

# Minimax Implementation (Dispatch)

```
def value(state):
```

    if the state is a terminal state: return the state's utility

    if the next agent is MAX: return max-value(state)

    if the next agent is MIN: return min-value(state)

```
def max-value(state):
```

    initialize  $v = -\infty$

    for each successor of state:

$v = \max(v, \text{min-value}(\text{successor}))$

    return  $v$

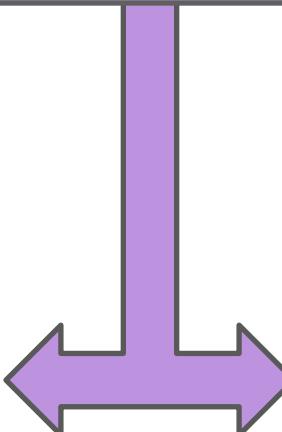
```
def min-value(state):
```

    initialize  $v = +\infty$

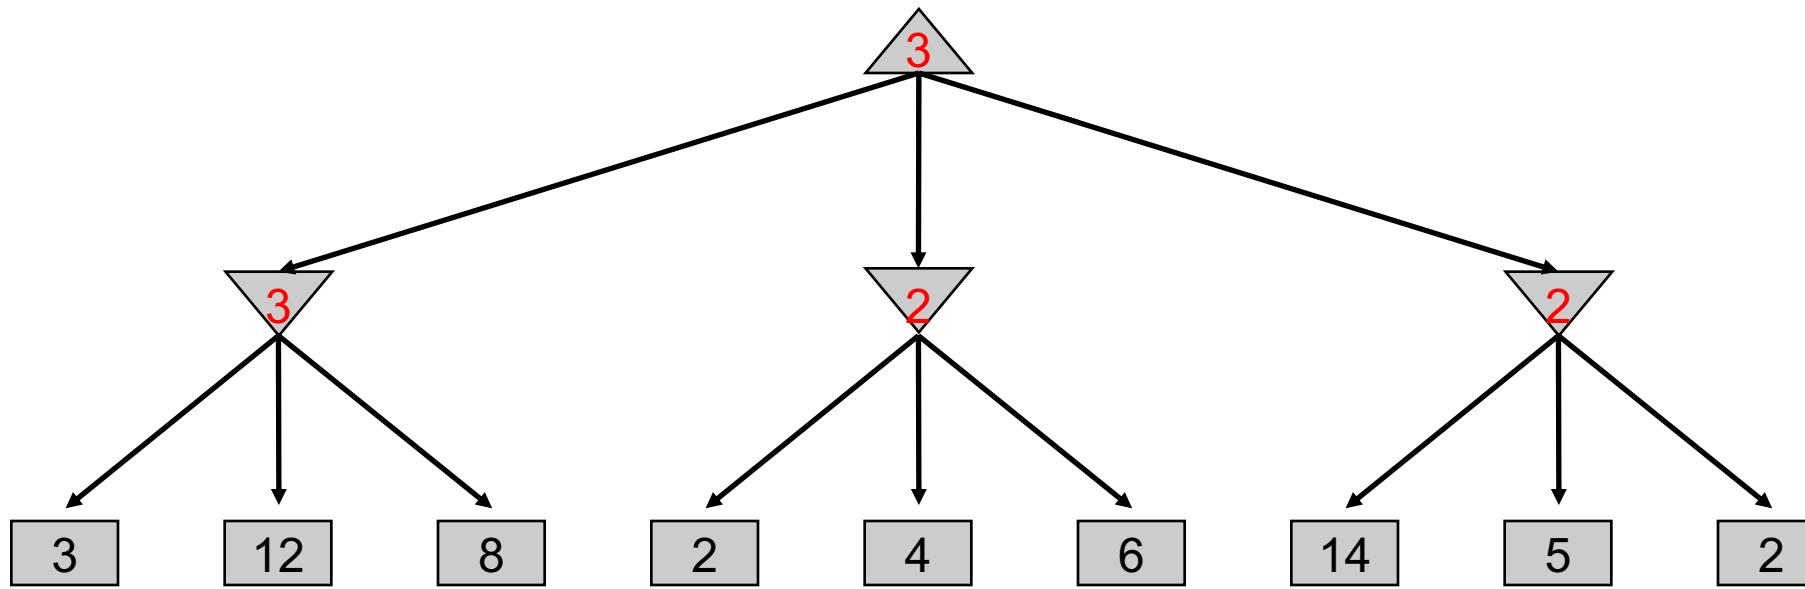
    for each successor of state:

$v = \min(v, \text{max-value}(\text{successor}))$

    return  $v$



# Minimax Example



# Minimax Efficiency

---

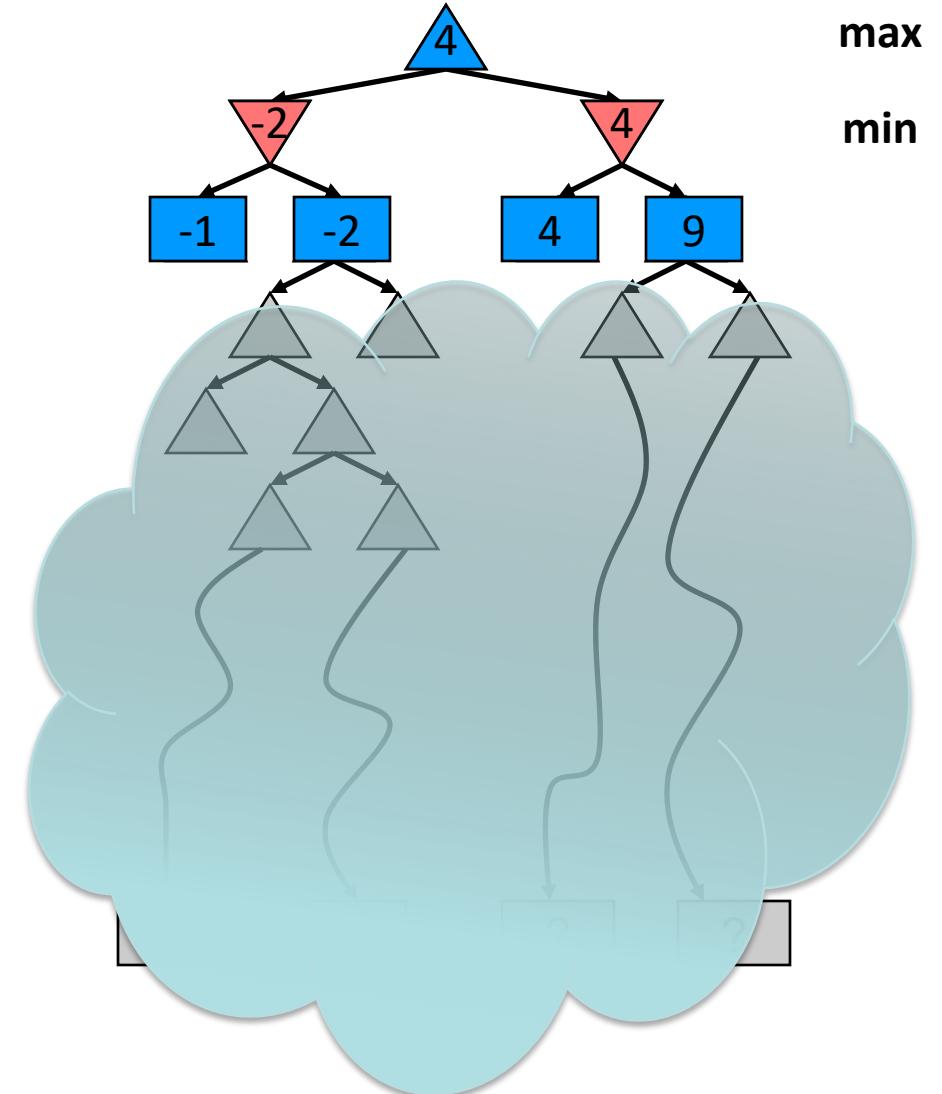
- How efficient is minimax?
  - Just like (exhaustive) DFS
  - Time:  $O(b^m)$
  - Space:  $O(bm)$
- Example: For chess,  $b \approx 35$ ,  $m \approx 100$ 
  - Exact solution is completely infeasible
  - But, do we need to explore the whole tree?

# Resource Limits

---

# Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution 1: Depth Limited Search
  - Search only to a preset **depth limit** or **horizon**
  - Use an **evaluation function** for non-terminal positions
- Guarantee of optimal play is gone
- Deeper search makes a BIG difference
- Example:
  - Suppose we have 100 seconds, can explore 10K nodes / sec
  - So can check 1M nodes per move
  - For chess, b=~35 so reaches about depth 4 – not so good

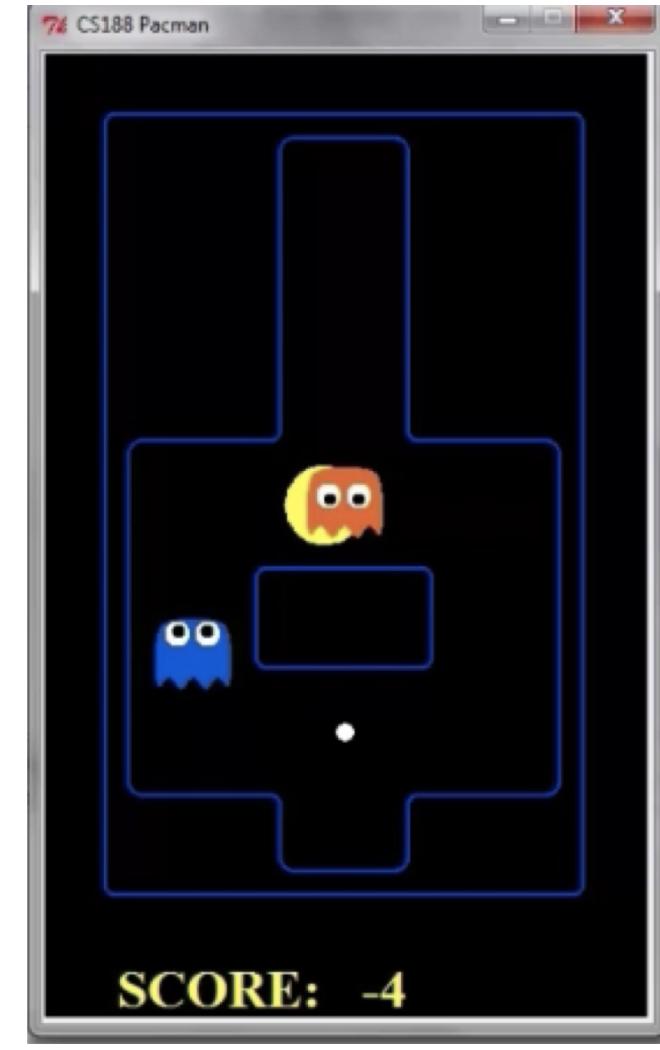


# Depth Matters

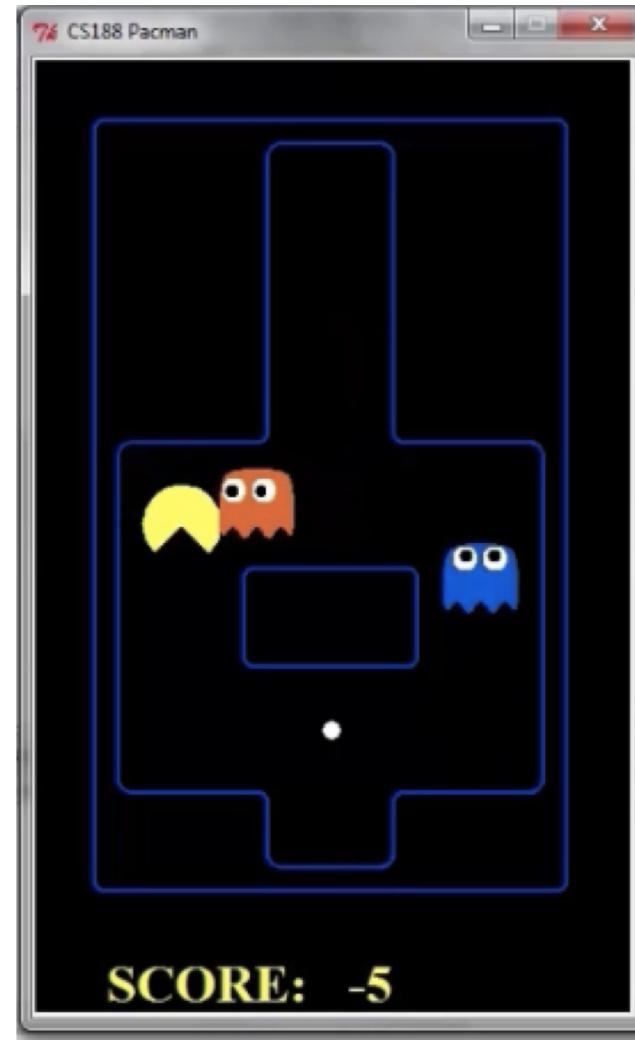
---

- Evaluation functions are always imperfect
- Deeper search => better play (usually)
- Or, deeper search gives same quality of play with a less accurate evaluation function

# Video of Demo Limited Depth (2)



# Video of Demo Limited Depth (10)

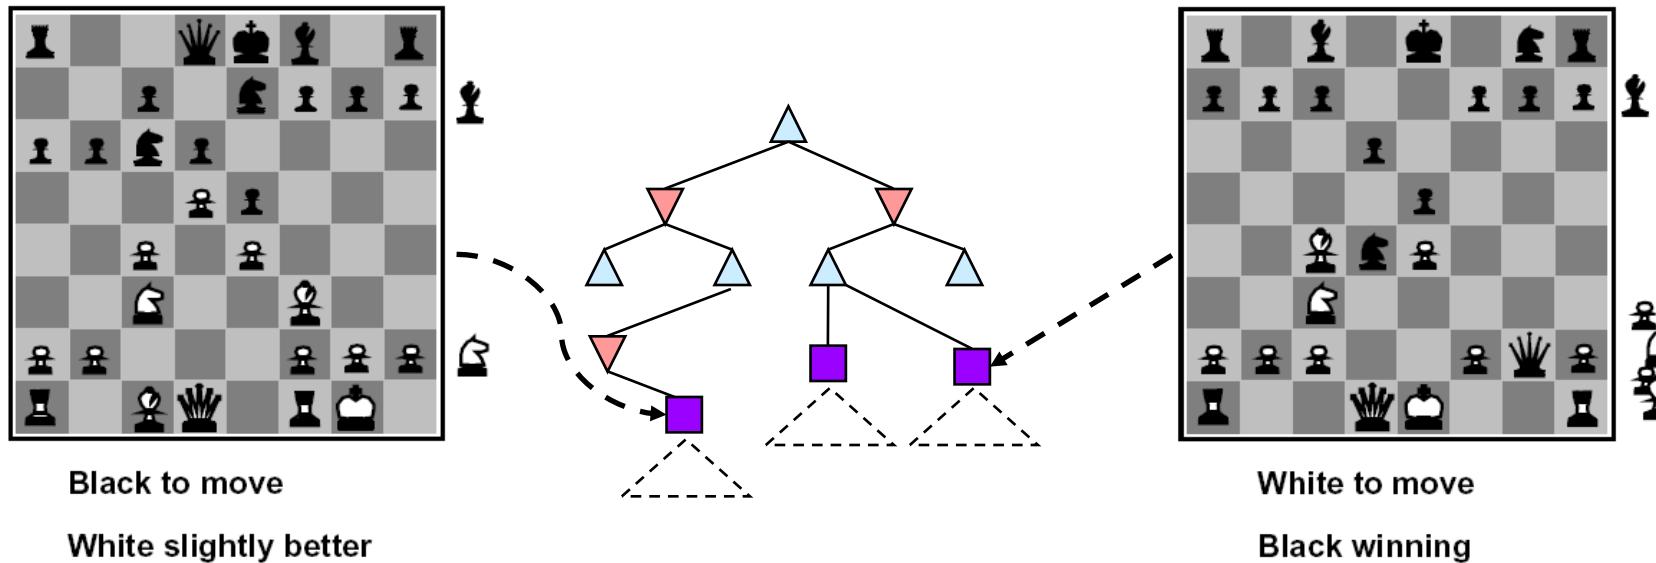


# Evaluation Functions

---

# Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



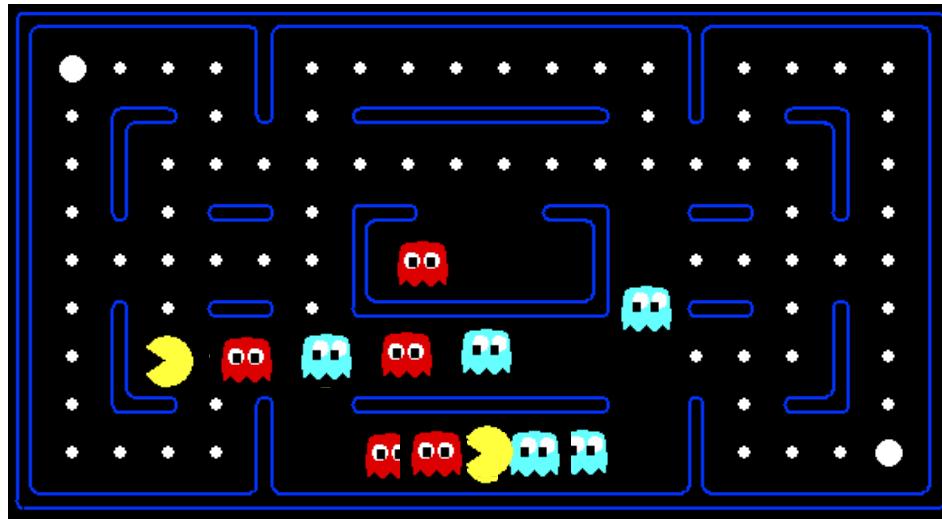
- Ideal evaluation function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = (\text{num white queens} - \text{num black queens})$ ,  
 $f_2(s) = (\text{num white pawns} - \text{num black pawns})$ ,

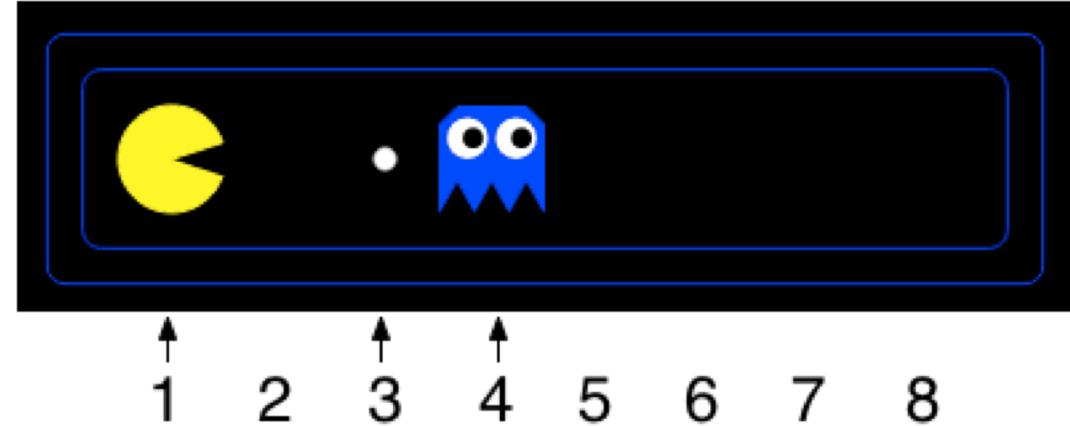
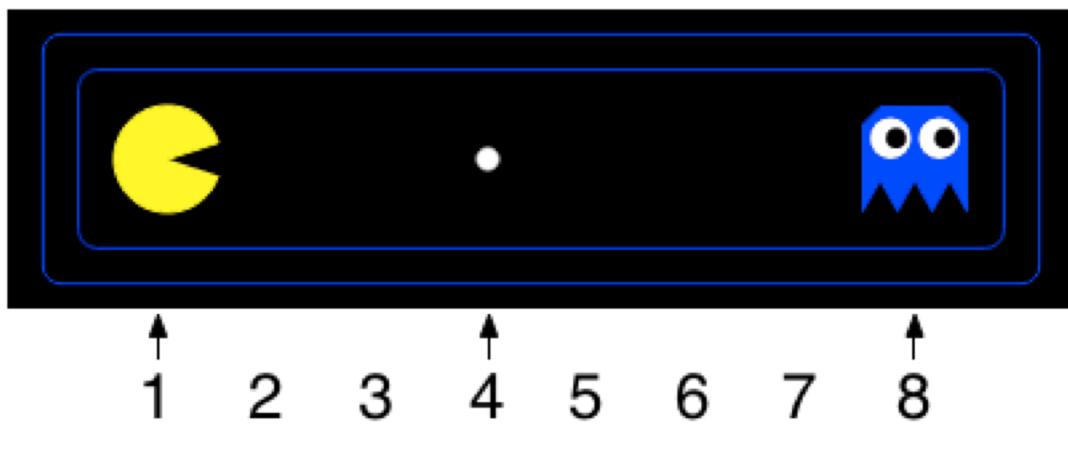
# Evaluation for Pacman

---



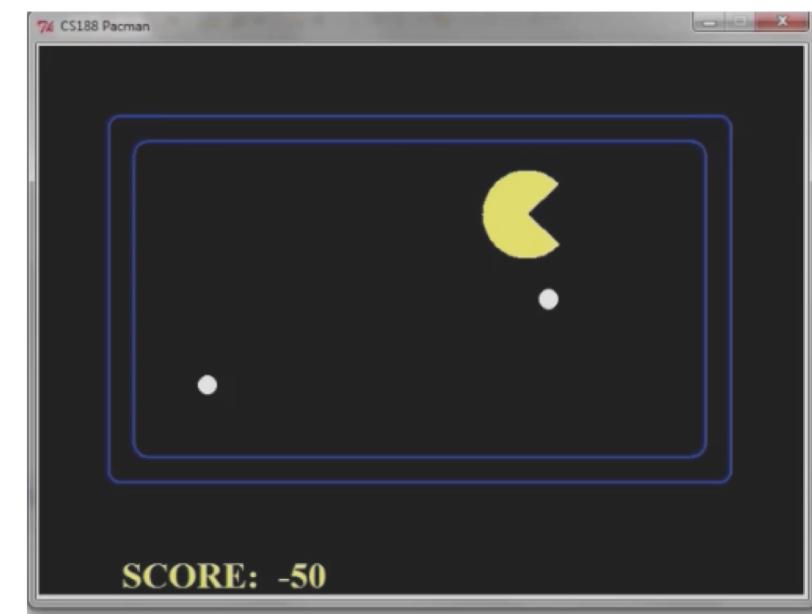
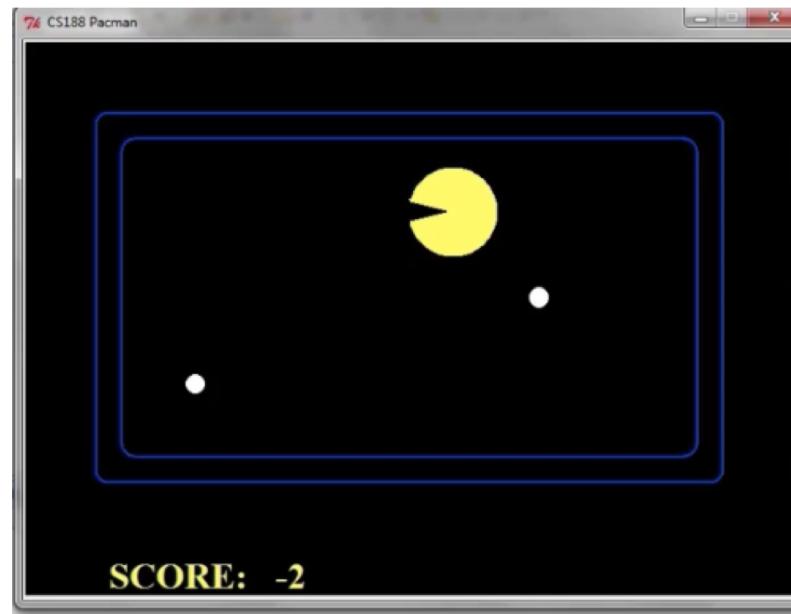
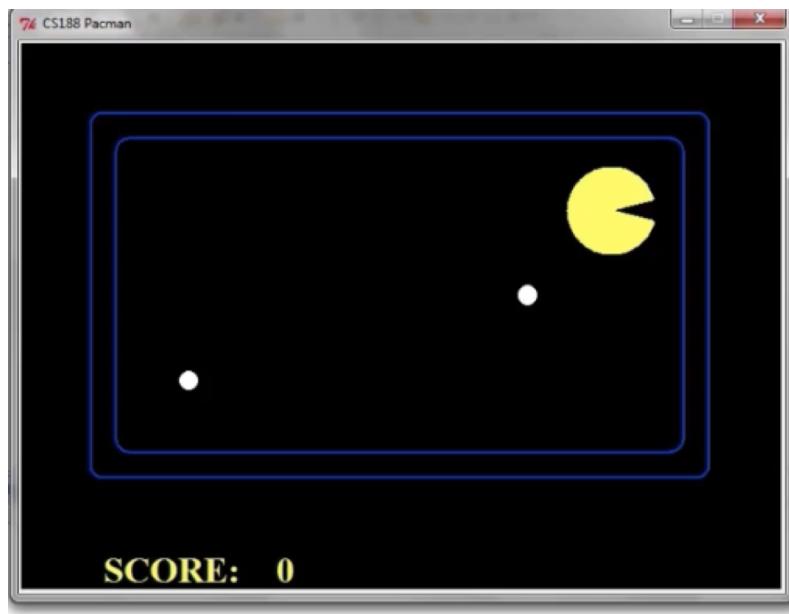
# Evaluation Function Quiz

For the two situations shown below, which evaluation functions will give the situation on the left a higher score than the situation on the right?



1.  $1 / (\text{Pac-Man's distance to the nearest food pellet})$
2. Pac-Man's distance to the nearest ghost
3. Pac-Man's distance to the nearest ghost +  $1 / (\text{Pac-Man's distance to the nearest food pellet})$
4. Pac-Man's distance to the nearest ghost +  $1000 / (\text{Pac-Man's distance to the nearest food pellet})$

# Video of Demo Thrashing (d=3)



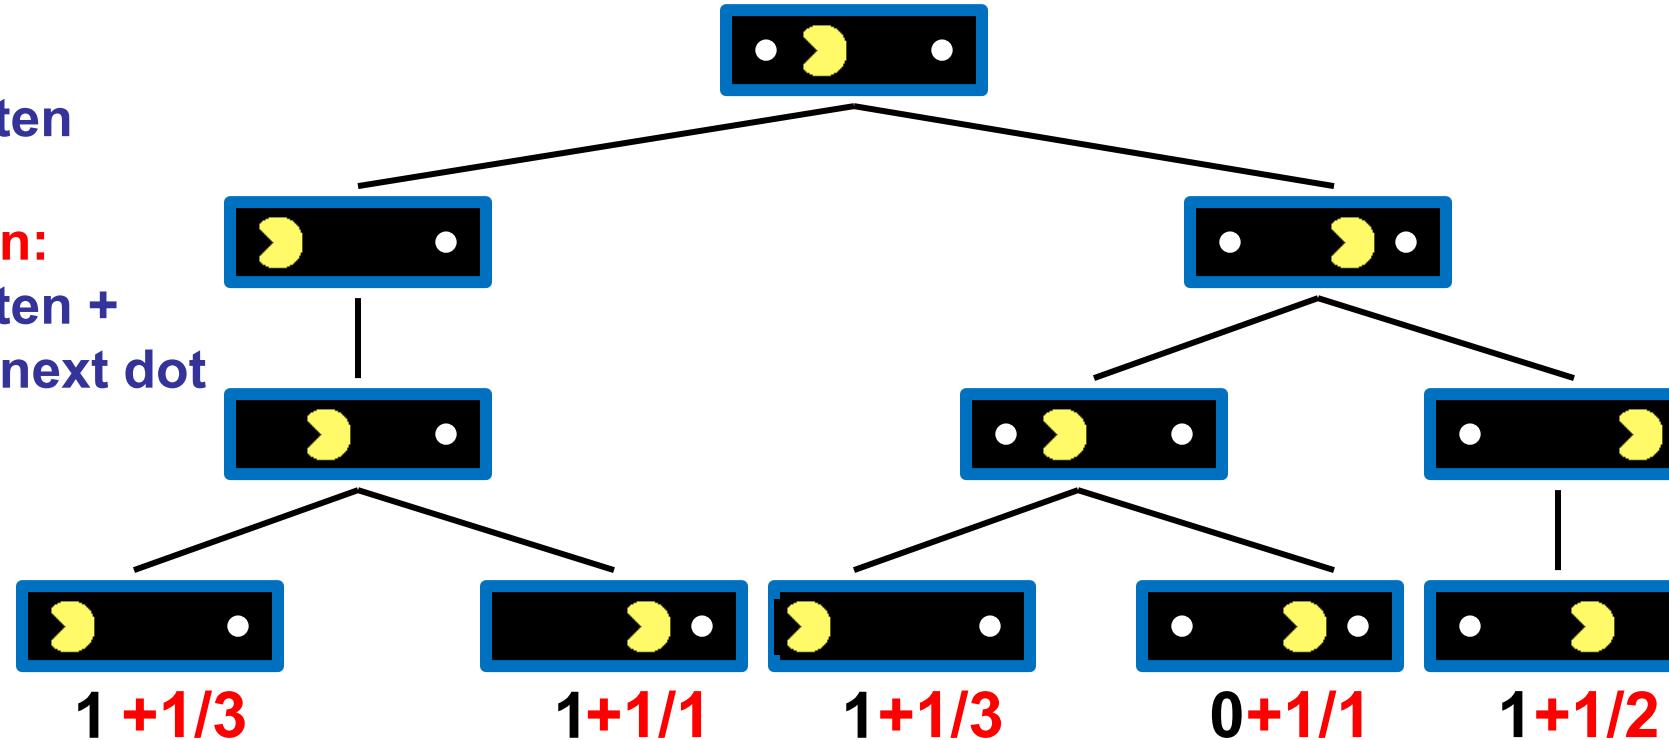
# Why Pacman Starves

**Eval Function:**

Number of Dots Eaten

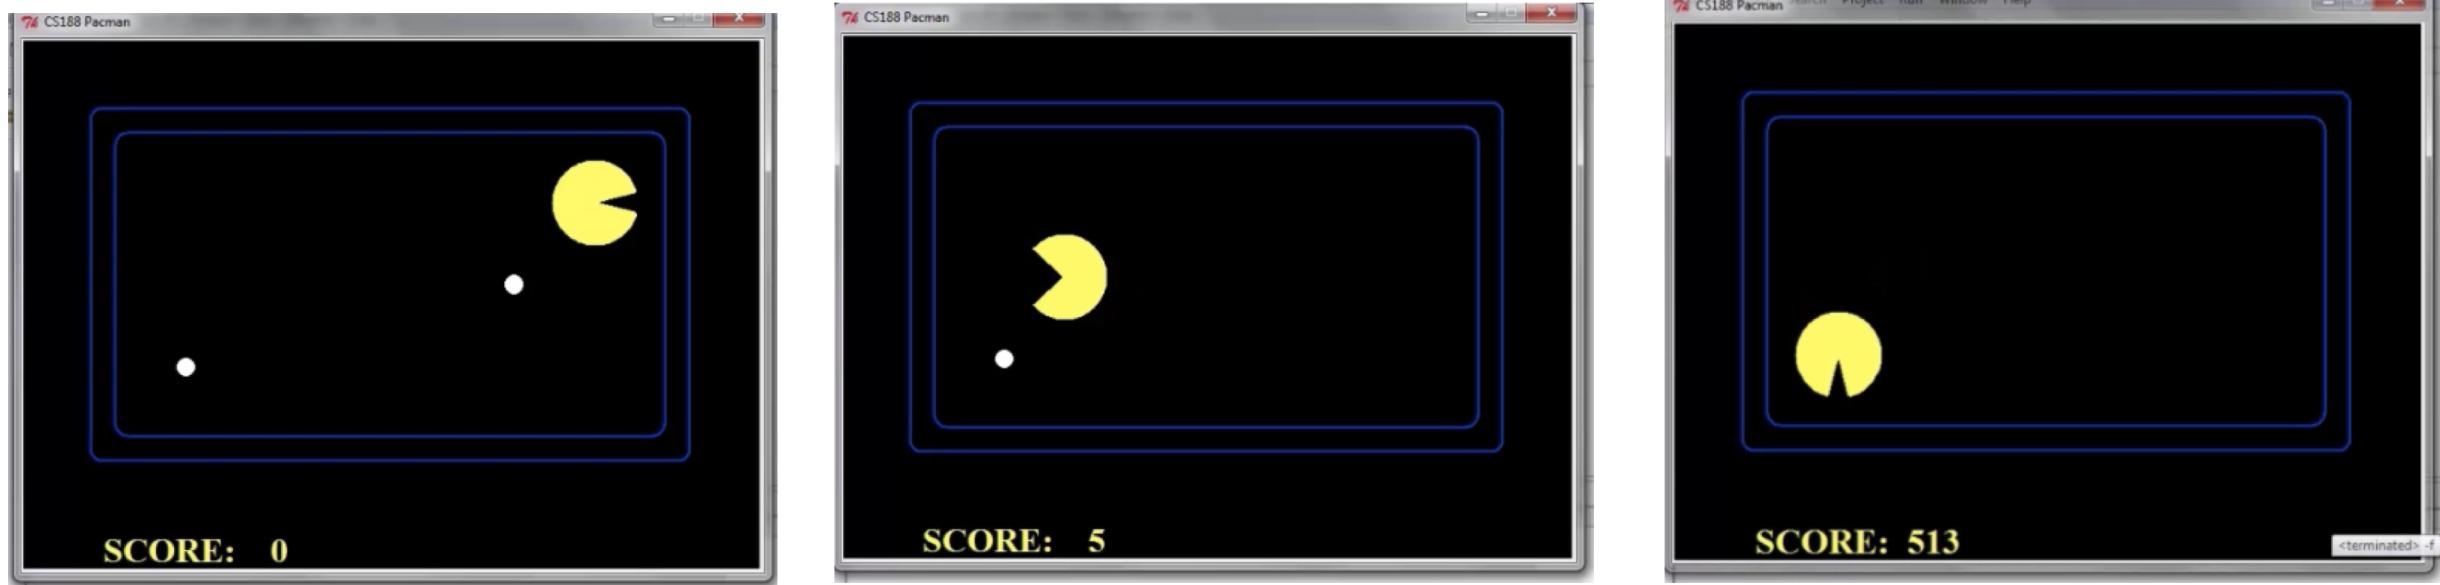
**Better Eval Function:**

Number of Dots Eaten +  
1/Closeness to the next dot

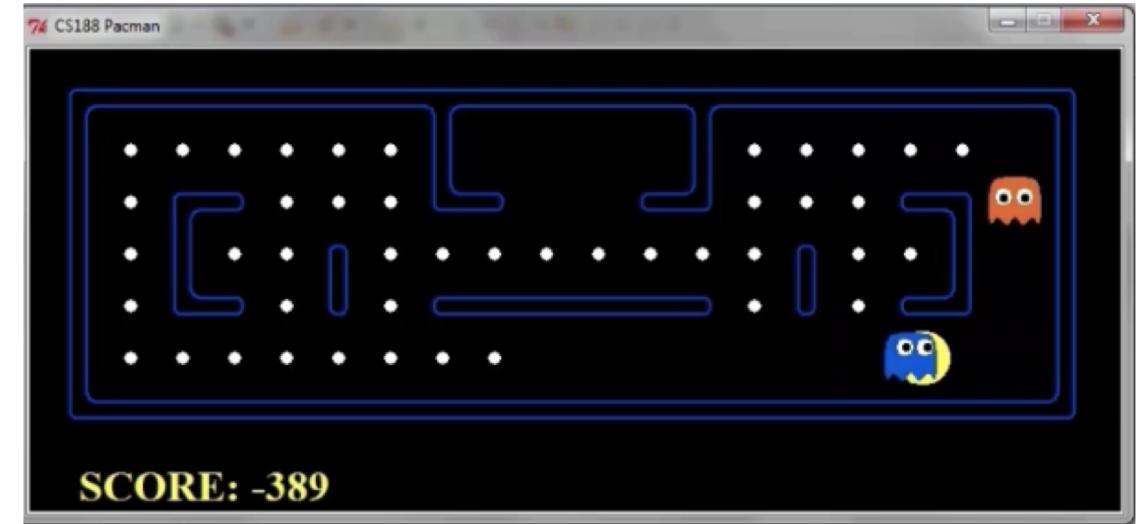
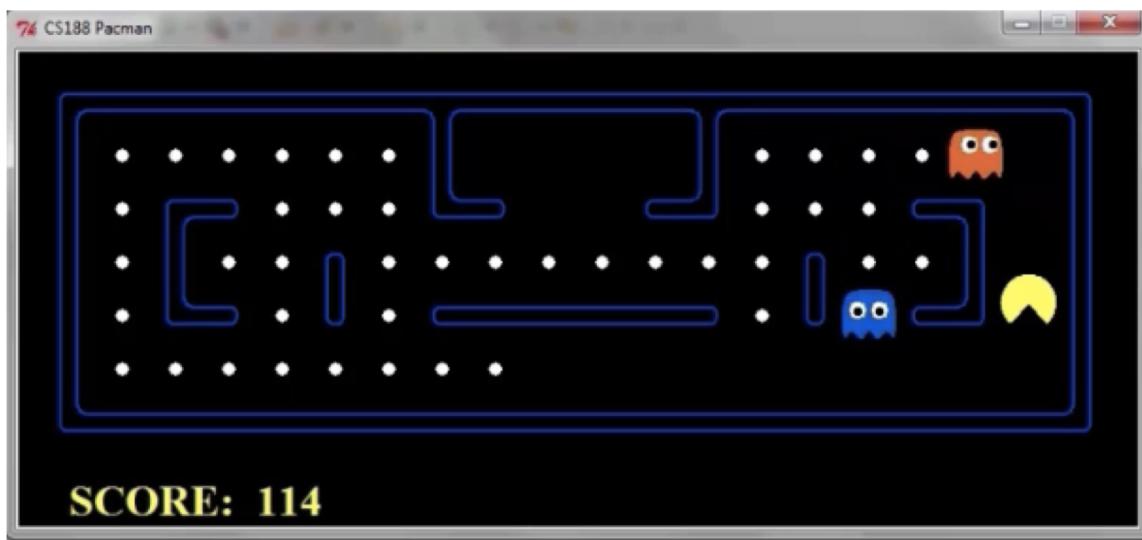
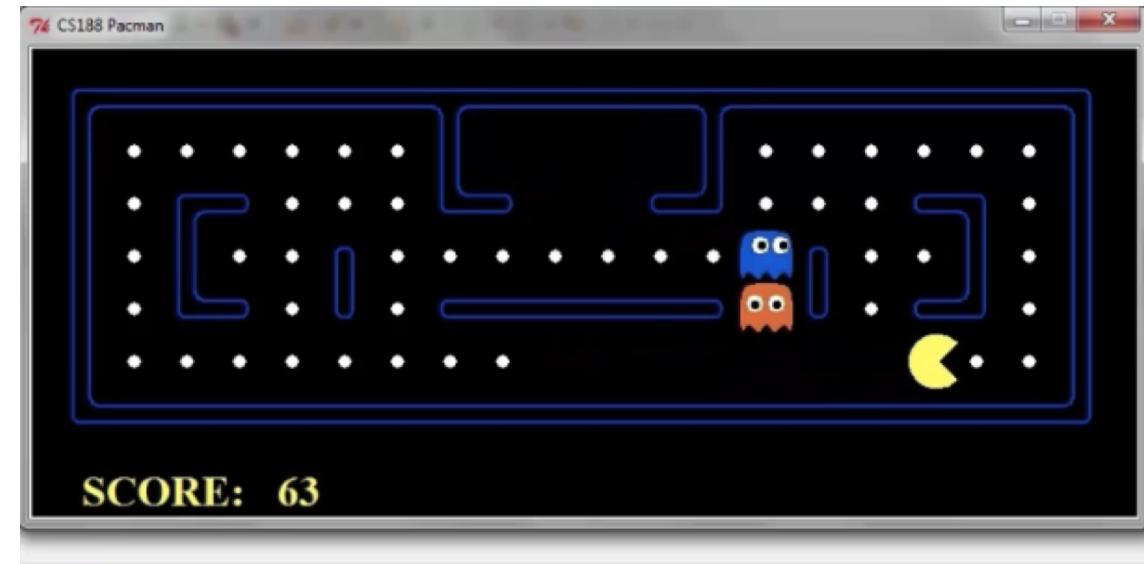
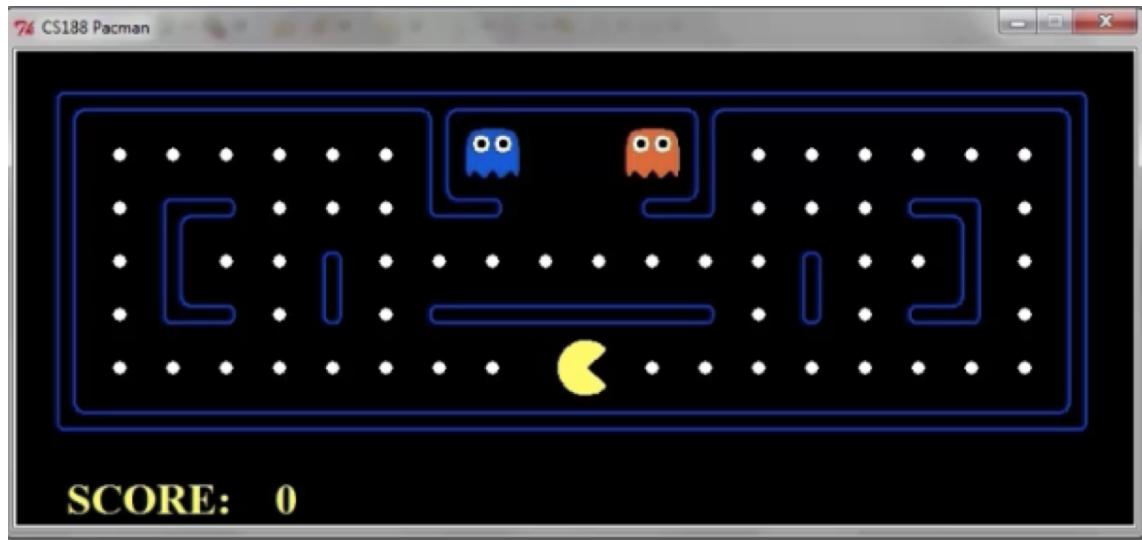


- A danger of replanning agents!
  - He knows his score will go up by eating the dot now
  - He knows his score will go up just as much by eating the dot later
  - Therefore, waiting seems just as good as eating.

# Video of Demo Thrashing -- Fixed (d=3)

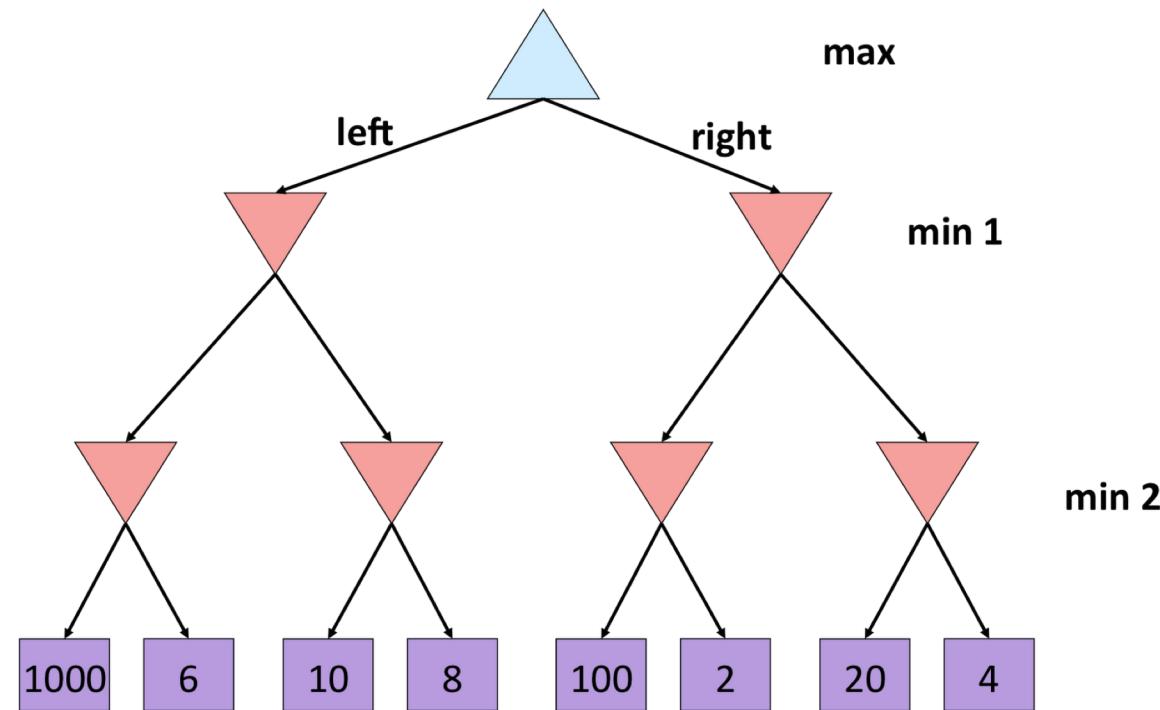


# Video of Demo Smart Ghosts (Coordination)



# Collaboration Quiz

Below is an example of a game tree with two minimizer players (min 1 and min 2), and one maximizer player.



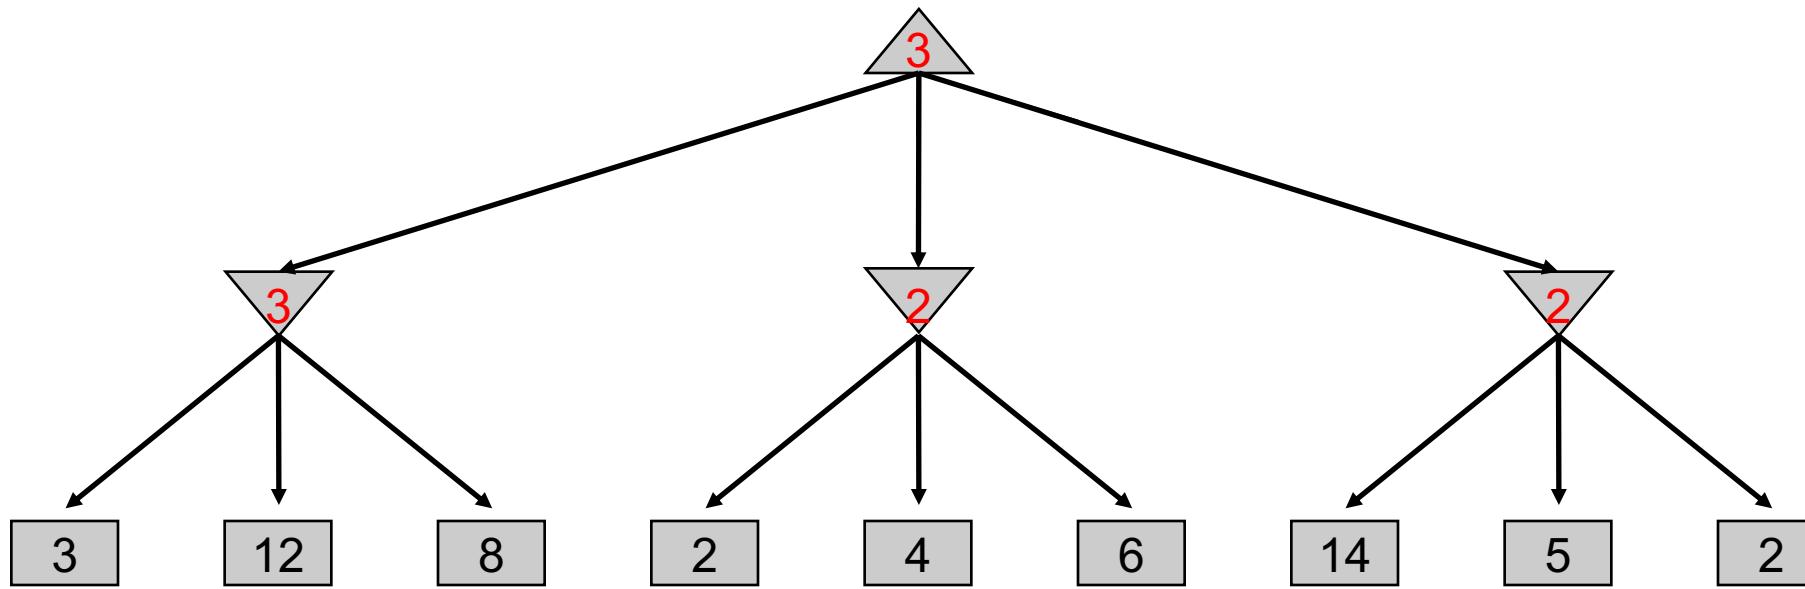
What is the minimax value of this game tree? 6

Which action will the maximizer take when playing according to the minimax strategy? Left

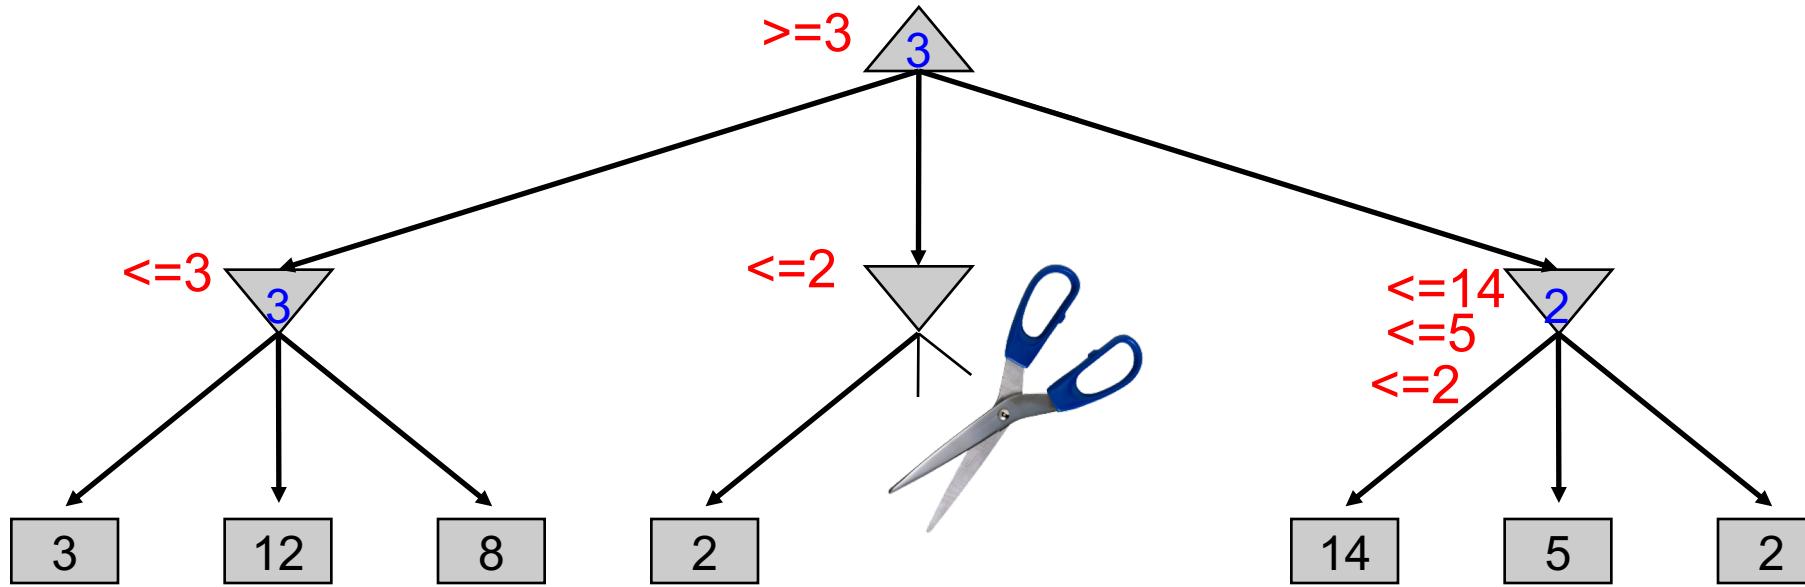
# Game Tree Pruning

---

# Minimax Example

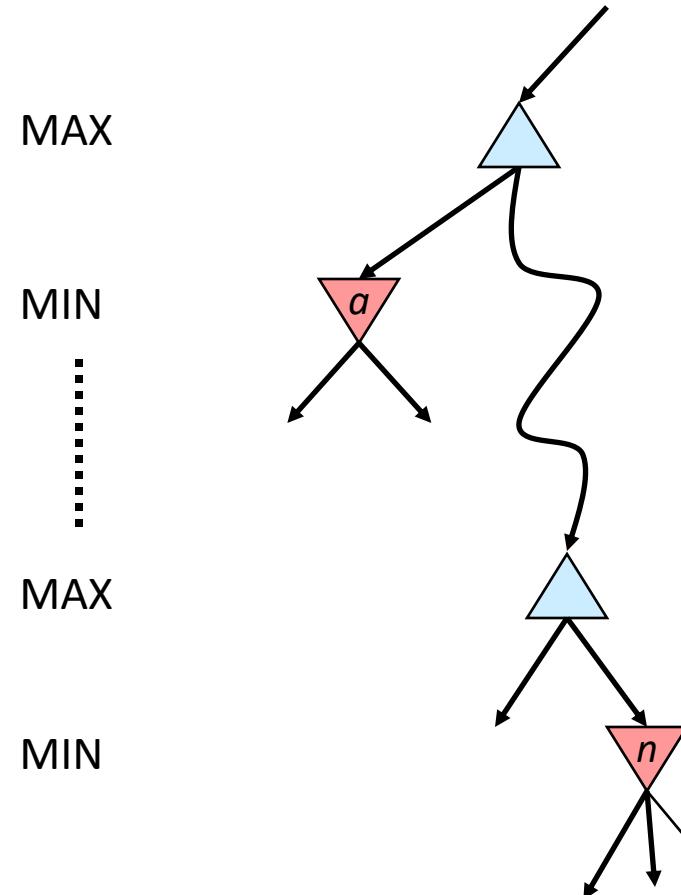


# Minimax Pruning



# Alpha-Beta Pruning

- General case (pruning children of MIN node)
  - We're computing the **MIN-VALUE** at some node  $n$
  - We're looping over  $n$ 's children
  - $n$ 's estimate of the children's min is dropping
  - Who cares about  $n$ 's value? **MAX**
  - Let  $\alpha$  be the best value that **MAX** can get so far at any choice point along the current path from the root
  - If  $n$  becomes worse than  $\alpha$ , **MAX** will avoid it, so we can prune  $n$ 's other children (it's already bad enough that it won't be played)
- Pruning children of **MAX** node is symmetric
  - Let  $\beta$  be the best value that **MIN** can get so far at any choice point along the current path from the root



# Alpha-Beta Implementation

$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

```
def max-value(state, α, β):  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, value(successor, α, β))  
        if v ≥ β  
            return v  
    α = max(α, v)  
    return v
```

```
def min-value(state , α, β):  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, value(successor, α, β))  
        if v ≤ α  
            return v  
    β = min(β, v)  
    return v
```

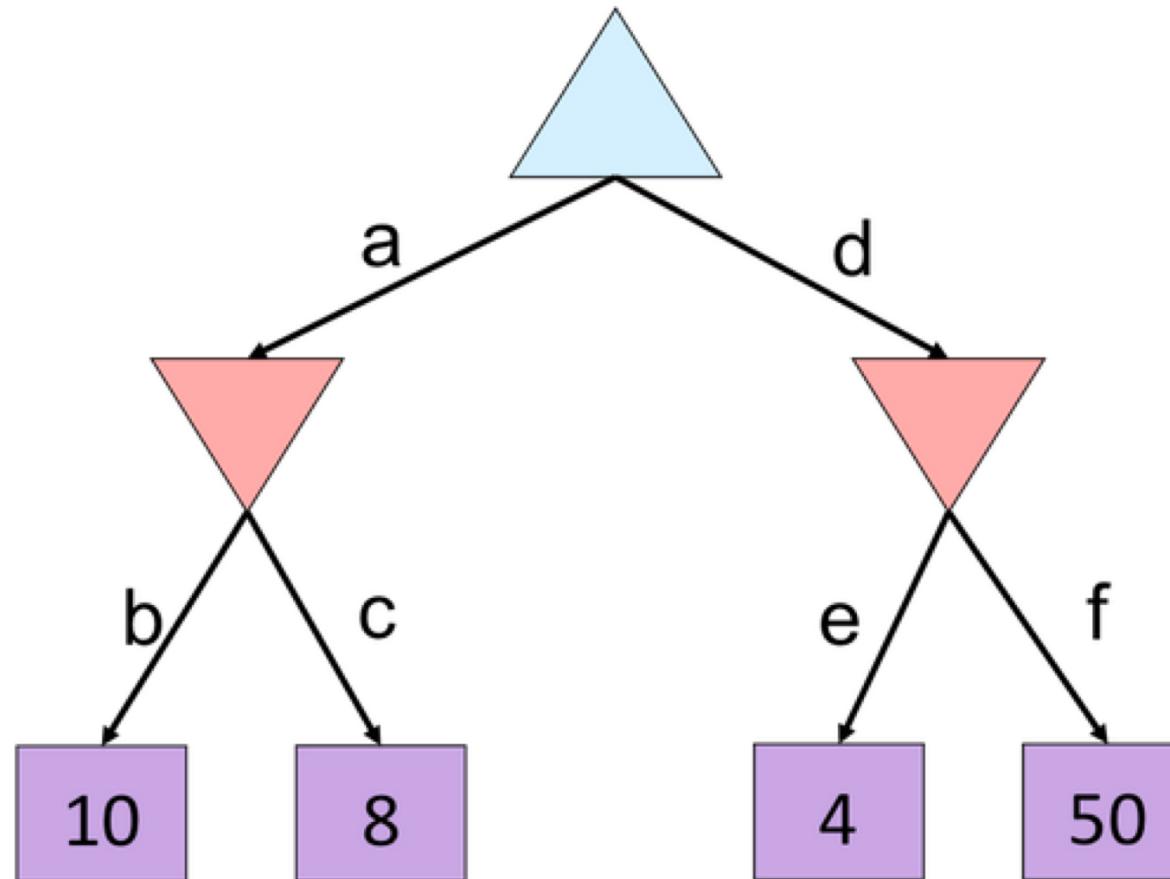
# Alpha-Beta Pruning Properties

---

- Theorem: This pruning has **no effect** on minimax value computed for the root!
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
  - Time complexity drops to  $O(b^{m/2})$
  - Doubles solvable depth!
  - 1M nodes/move => depth=8, respectable

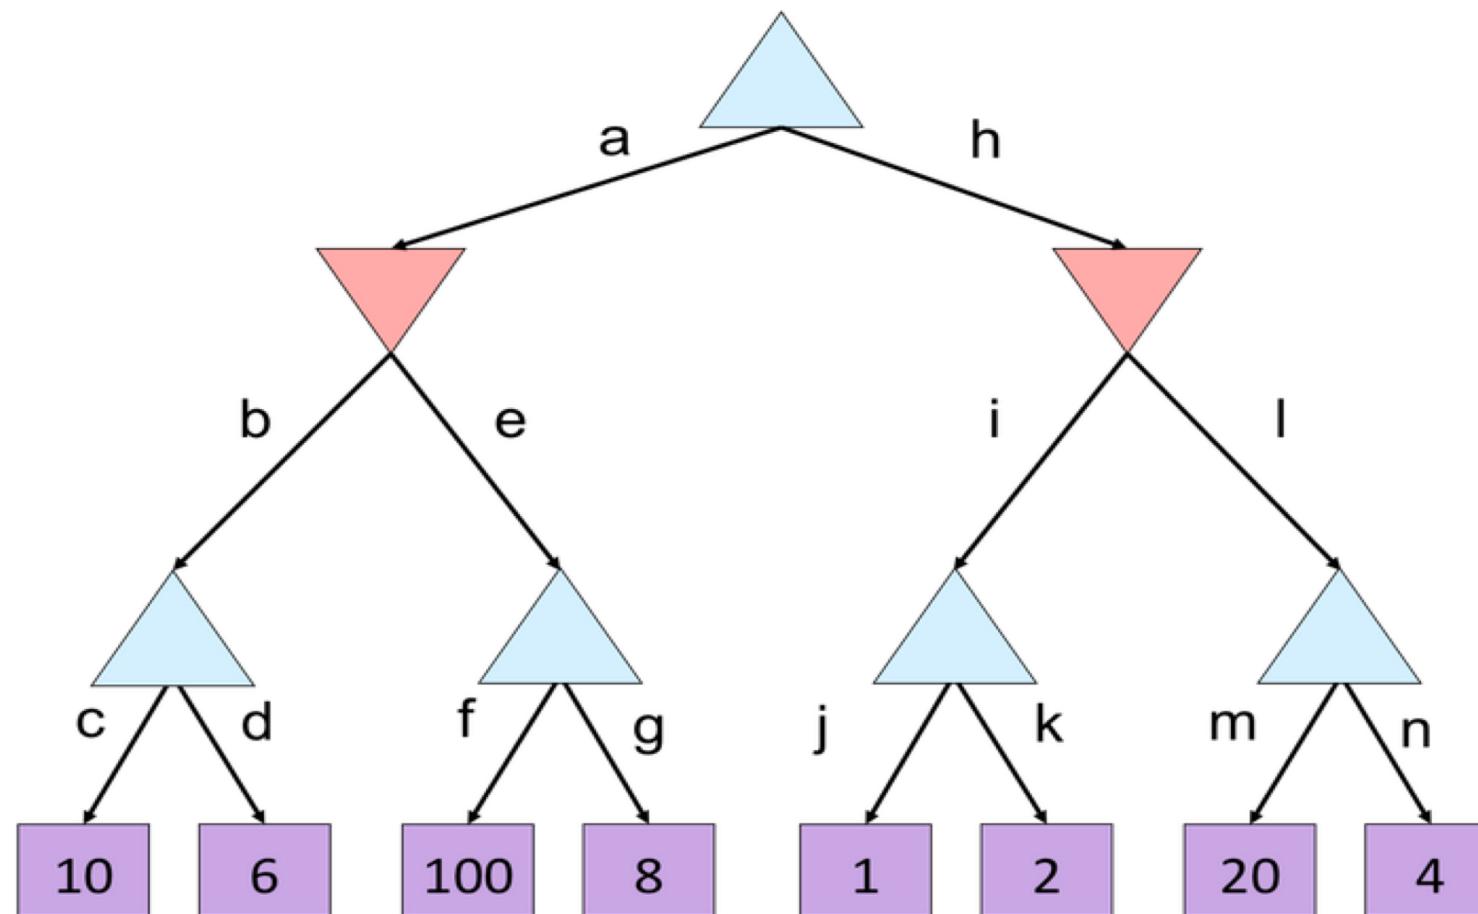
# Alpha-Beta Quiz

For the game tree shown below, which branches will be pruned by alpha-beta pruning? f



# Alpha-Beta Quiz 2

For the game tree shown below, which branches will be pruned by alpha-beta pruning? **g and l**



# Reading

---

- Chapter 5.1-5.5 in the AIMA textbook